

# MVE137

## Probability and Statistical Learning Using Python

### Supervised Learning

Alexandre Graell i Amat

[alexandre.graell@chalmers.se](mailto:alexandre.graell@chalmers.se)

<https://sites.google.com/site/agraellamat>

September 28, 2021



**CHALMERS**

# Supervised learning

## Statistical learning

**Goal:** Infer a **predictive function** (**model**), such that it can be used to predict the output for new, yet unseen data.

Examples: predicting housing or stock prices, image classification, ...

## Supervised learning

**Goal:** Given a training set  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ , identify an algorithm to predict the outcome  $y$  for a new (yet unseen) data point  $\mathbf{x}$ .

- Learning a model from labeled data
  - Predicting output of new data based on the learned model
- 
- $\mathbf{x}_i$ : free variables (features, predictors, covariates, domain points)
  - $y_i$ : target variables (dependent variables, labels, responses)

# Supervised learning

Three classes of responses:

- **Continuous** (**quantitative**): take on numerical values
- **Discrete** (**qualitative**, **categorical**): a discrete set of categories
- **Order categorical**: the order is important

Two learning algorithms:

- **Regression**: predict a quantitative output
- **Classification**: predict a qualitative output

# Statistical learning and decision theory

**Goal:** Given a training set  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ , predict  $y$  for a new (yet unseen) data point  $\mathbf{x}$ .

Impossible if no information on the mechanism relating  $\mathbf{x}$  and  $y$ !

We may assume that  $\mathbf{x}$  and  $y$  are related via a function

$$y = \tilde{f}(\mathbf{x})$$

**Goal:** Find best possible approximation of  $\tilde{f}(\mathbf{x})$ ,  $f(\mathbf{x})$  (prediction model).  
Predict the outcome  $y$  for  $\mathbf{x}$  as  $\hat{y} = f(\mathbf{x})$ .

- Treat  $\mathbf{x}$  and  $y$  as random variables, and

$$(\mathbf{x}_i, y_i) \sim_{i.i.d.} p(\mathbf{x}, y), \quad i \in [N].$$

# Statistical learning and decision theory

How should we choose  $f(\mathbf{x})$ ?

- Loss function  $\ell(y, \hat{y}) = \ell(y, f(\mathbf{x}))$ : cost (loss or risk) incurred when the correct value is  $y$  while the estimate is  $\hat{y}$
- Quadratic loss function:

$$\ell(y, \hat{y}) = (y - \hat{y})^2 = (y - f(\mathbf{x}))^2,$$

Expected prediction error (expected generalization loss/error):

$$\begin{aligned} L(\hat{y}) &= \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] \\ &= \int \int \ell(y, f(\mathbf{x})) p(\mathbf{x}, y) d\mathbf{x} dy. \end{aligned}$$

# Statistical learning and decision theory

How should we choose  $f(\mathbf{x})$ ?

Optimal prediction  $\hat{y}(\mathbf{x})$  obtained by minimizing generalization loss:

$$\begin{aligned} f^*(\mathbf{x}) &= \operatorname{argmin}_f L(\hat{y}) \\ &= \operatorname{argmin}_f \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] \\ &= \operatorname{argmin}_f \mathbb{E}_{\mathbf{x} \sim p_{\mathbf{x}}} [\mathbb{E}_{y \sim p_{y|\mathbf{x}}} [\ell(y, f(\mathbf{x}))]] \end{aligned}$$

Hence, it suffices to solve

$$f^*(\mathbf{x}) = \operatorname{argmin}_f \mathbb{E}_{y \sim p_{y|\mathbf{x}}} [\ell(y, f(\mathbf{x}))]$$

# Statistical learning and decision theory

For the quadratic loss  $\ell(y, \hat{y}) = (y - \hat{y})^2$ ,

$$f^*(\mathbf{x}) = \operatorname{argmin}_f \mathbb{E}_{y \sim p_{y|\mathbf{x}}}[\ell(y, f(\mathbf{x}))] =$$

# Linear regression

**Linear regression:** Assumes linear model for  $f(\mathbf{x})$ ,

$$\hat{y} = f(\mathbf{x}) = \beta_0 + \sum_{j=1}^p x_j \beta_j,$$

with  $\mathbf{x} = (x_1, x_2, \dots, x_p)^\top$ .

$\beta_0$ : intercept or bias

- For simplicity, we will write

$$\tilde{\mathbf{x}} = (1, x_1, \dots, x_p)^\top \quad \text{and} \quad \boldsymbol{\beta} = (\beta_0, \dots, \beta_p)^\top$$

so that

$$\hat{y} = \sum_{j=0}^p \tilde{x}_j \beta_j = \tilde{\mathbf{x}}^\top \boldsymbol{\beta}$$

**Linear regression** assumes  $\tilde{f}(\mathbf{x}) \approx \mathbf{x}^\top \boldsymbol{\beta}$  or, equivalently,  $\mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] \approx \tilde{\mathbf{x}}^\top \boldsymbol{\beta}^*$ .



## Linear regression: Optimal $\beta$

$$\begin{aligned}\beta^* &= \operatorname{argmin}_{\beta} L(\hat{y}) \\ &= \operatorname{argmin}_{\beta} \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] \\ &= \operatorname{argmin}_{\beta} \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [(y - \mathbf{x}^\top \beta)^2]\end{aligned}$$

**Idea:** Approximate the expectation by the **empirical average** over the  $N$  training points  $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^N$ ,

$$\begin{aligned}\beta^* &= \operatorname{argmin}_{\beta} \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \beta)^2 \\ &= \operatorname{argmin}_{\beta} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \beta)^2\end{aligned}$$

## Least squares linear regression

$$\begin{aligned}\beta^* &= \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \beta)^2 \\ &= \underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|^2 \\ &= \underset{\beta}{\operatorname{argmin}} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)\end{aligned}$$

with

$$\mathbf{X} = \begin{pmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,p} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,p} \\ \vdots & \vdots & \vdots & \dots & \vdots \\ 1 & x_{N,1} & x_{N,2} & \dots & x_{N,p} \end{pmatrix}$$

and  $\mathbf{y} = (y_1, \dots, y_N)^\top$

## Least squares linear regression

$$\beta^* = \underset{\beta}{\operatorname{argmin}} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)$$

# Least squares linear regression for classification

Binary classification:  $y \in \{a, b\}$



**Idea:** Encode  $y = a \implies y = 0$  and  $y = b \implies y = 1$  and apply plain linear regression plus **thresholding**:

$$\hat{y}(\mathbf{x}) = \begin{cases} 0 & \text{if } f^*(\mathbf{x}) = \mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] = \mathbf{x}^\top \boldsymbol{\beta}^* \leq 0.5 \\ 1 & \text{otherwise} \end{cases}$$

**Observations:**

- The linear model determines the decision boundary  $\{\mathbf{x} : \mathbf{x}^\top \boldsymbol{\beta}^* = 0.5\}$
- We can interpret  $f^*(\mathbf{x}) = \mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] = \mathbf{x}^\top \boldsymbol{\beta}^*$  as the probability  $p(y = 1|\mathbf{x})$

## Least squares linear regression for classification



- 100 samples from two bivariate Gaussian distributions
- Blue:  $\mathcal{N}((1, 0)^T, \mathbf{I})$
- Orange:  $\mathcal{N}((0, 1)^T, \mathbf{I})$

**Goal:** For a new coordinate  $\mathbf{x} = (x_1, x_2)$ , determine to which of the Gaussians corresponds to.

## $k$ -Nearest neighbor regression

$k$ -Nearest neighbor regression: Given  $\mathbf{x}$ , it predicts  $y$  as

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in \mathcal{N}_k(\mathbf{x})} y_i(\mathbf{x}_i)$$

$\mathcal{N}_k(\mathbf{x})$ : set of  $k$  nearest neighbors to  $\mathbf{x}$  in the training set.

- For  $\mathbf{x} \in \mathbb{R}^p$ , we consider the Euclidean distance  $d(\mathbf{x}, \mathbf{x}_i) = \|\mathbf{x} - \mathbf{x}_i\|^2 \longrightarrow \mathcal{N}_k(\mathbf{x})$  is the set  $\{\mathbf{x}_i\}$  closest (in Euclidean distance) to  $\mathbf{x}$ .

## $k$ -Nearest neighbor regression

Optimal solution:

$$\hat{y}(x) = f^*(x) = \mathbb{E}_{y|x}[y|x]$$

$k$ -nearest neighbors tries to accomplish this directly!

1. Replacing  $\mathbf{x} = x$  by neighborhood of  $x$  in the training data  $\mathcal{N}_k(x)$
2. Replacing expectation by average over the  $k$  training neighbors

We make **two approximations**:

- Expectation approximated by sample average
- Conditioning at point  $x$  relaxed to conditioning on region close  $x$

## $k$ -Nearest neighbor regression

What's the role of  $k$ ?

For a larger  $k$  ...

- Average is more accurate and stable (reduced variance)
- Neighborhood is bigger and less representative of  $\mathbf{x} = \mathbf{x}$  (increased bias)

Under mild regularity conditions on  $p(\mathbf{x}, y)$ , as  $N \rightarrow \infty$ ,  $k(N) \rightarrow \infty$  and  $k(N)/N \rightarrow 0$ ,  $\hat{y}(\mathbf{x}) \rightarrow \tilde{f}(\mathbf{x})$ ,  $\forall \mathbf{x}$ .





## $k$ -Nearest neighbors for classification

**Idea:** Replace averaging by a **majority vote**,

$$\hat{y}(\mathbf{x}) = \mathbb{1} \left\{ \frac{1}{k} \sum_{\mathbf{x}_i \in \mathcal{N}_k(\mathbf{x})} \mathbb{1}\{y_i(\mathbf{x}_i) = 1\} > 0.5 \right\}$$

1. Find  $\mathcal{N}_k(\mathbf{x})$
2. **Majority vote:** Assign  $\mathbf{x}$  to the class that most predictors in  $\mathcal{N}_k(\mathbf{x})$  belong

## $k$ -Nearest neighbors for classification



- Left: 15-nearest neighbors
- Right: 1-nearest neighbor

Error on training data **decreases** with **decreasing**  $k$ , and is zero for  $k = 1$ .

Is  $k = 1$  **optimal**?

# Least squares regression vs $k$ -nearest neighbors regression

## Least squares regression:

- Assumes  $\tilde{f}(x)$  well approximated by a globally linear function
- Very smooth boundary
- Stable to fit
- Linear decision boundary (strong assumption!)
- Low variance and (potentially) high bias

## $k$ -nearest neighbors regression:

- No stringent assumptions about underlying data
- Can adapt to any shape of the data
- Not smooth boundary (for small  $k$ )
- Unstable to fit (for small  $k$ )
- high variance and low bias

# Parametric vs nonparametric models

**Parametric models:** Build  $f(x)$  as a parametric model that applies to the **whole space**.

1. Select parametric model (hypothesis class), with fixed number of parameters
2. Learn parameters to fit the training data  $\mathcal{D}$

**Nonparametric models:** Don't make explicit assumptions about form  $f(x)$ , but describe it in terms of **local behavior** of the training data in the region near  $x$ .

1. Seek an estimate of  $f$  that gets as close to the data points as possible without being too wiggly
2. Advantage: accurately fit a wider range of possible shapes for  $f$
3. Disadvantage: number parameters grows with amount of training data

# Classification and statistical learning

- $y \in \mathcal{C} = \{C_1, \dots, C_K\}$ ,  $|\mathcal{C}| = K$
- We assume that  $\mathbf{x}$  and  $y$  are related via function  $\tilde{f}(\mathbf{x})$  (rule)

**Goal:** learn a rule  $f(\mathbf{x})$  which maps  $\mathbf{x}$  to one of the classes  $\{C_1, \dots, C_K\}$ .

How should we choose  $f(\mathbf{x})$ ?

- 0-1 loss function:

$$\ell(y, f(\mathbf{x})) = \begin{cases} 0 & y = f(\mathbf{x}) \\ 1 & y \neq f(\mathbf{x}) \end{cases}$$

Equivalently,

$$\ell(y, f(\mathbf{x})) = \mathbb{1}\{y \neq f(\mathbf{x})\}$$

# Classification and statistical learning

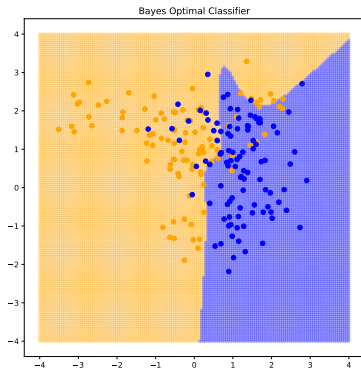
Optimal rule  $f^*(\mathbf{x})$ :

$$\begin{aligned} f^*(\mathbf{x}) &= \operatorname{argmin}_f \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [\ell(y, f(\mathbf{x}))] \\ &= \operatorname{argmin}_f \mathbb{E}_{\mathbf{x} \sim p_{\mathbf{x}}} \left[ \mathbb{E}_{y \sim p_{y|\mathbf{x}}} [\ell(y, f(\mathbf{x}))] \right] \\ &= \operatorname{argmin}_f \mathbb{E}_{\mathbf{x} \sim p_{\mathbf{x}}} \left[ \mathbb{E}_{y \sim p_{y|\mathbf{x}}} [\mathbb{1}\{y \neq f(\mathbf{x})\}] \right] \end{aligned}$$

# Classification and statistical learning

$$f^*(\mathbf{x}) = \operatorname{argmin}_f \mathbb{E}_{\mathbf{x} \sim p_{\mathbf{x}}} \left[ \mathbb{E}_{y \sim p_{y|\mathbf{x}}} [\mathbb{1}\{y \neq f(\mathbf{x})\}] \right]$$

# Classification and statistical learning





# Classification and statistical learning

Optimal classification: Bayes' classifier

$$f^*(\mathbf{x}) = \operatorname{argmax}_{i \in [K]} p(y = i | \mathbf{x})$$

$k$ -nearest neighbors: Approximates the optimal solution as

$$f^*(\mathbf{x}) = \operatorname{argmax}_{j \in [K]} \frac{1}{k} \sum_{\mathbf{x}_i \in \mathcal{N}_k(\mathbf{x})} \mathbb{1}\{y_i = j\}$$

# Curse of dimensionality (regression)

Optimal solution:


$$\hat{y}(x) = f^*(x) = \mathbb{E}_{y|x}[y|x]$$

$k$ -nearest neighbors: Approximates the optimal solution as

$$\hat{y}(x) = \frac{1}{k} \sum_{x_i \in \mathcal{N}_k(x)} y_i(x_i)$$

Can we accurately approximate the optimal solution by considering a very large training set? How large?

# Curse of dimensionality



## 3-nearest neighbors

- Blue and red: Training data points  $\mathbf{x}_i \in \mathcal{X} = [-1, 1]^2$  (10, 100, 1000)
- Black: New data point
- Red: 3 nearest neighbors

As  $N$  increases:

$$\hat{y}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_i \in \mathcal{N}_k(\mathbf{x})} y_i(\mathbf{x}_i) \rightarrow \mathbb{E}[y|\mathbf{x}]$$

# Curse of dimensionality


... but  $k$ -nearest neighbors (and other local methods) do not work well with **high-dimensional** inputs!

**Curse of dimensionality:** number of points exponential in number of dimensions!

1. Nearest neighbors not so close to  $x$
2.  $k$ -NNs of  $x$  closer to the boundary of  $\mathcal{X}$
3. Need a prohibitive number of training samples to densely sample  $\mathcal{X} \in \mathbb{R}^p$

(see “The elements of statistical learning,” Section 2.5, for 2 and 3)

For large  $p$ , the nearest neighbors are not so close




$k$ -nearest neighbors to data with training points uniformly distributed in a  $p$ -dimensional unit hypercube,  $\mathcal{X} = [0, 1]^p$ .

Want to estimate the **density** of class labels around a test point  $\mathbf{x}$  by growing a **hypercube** around  $\mathbf{x}$  until we capture a fraction  $\rho$  of the data points.

**Expected length** of the **side** of the smallest hypercube containing a fraction  $\rho$  of the data points:

For large  $p$ , the nearest neighbors are not so close



- Estimate based on 10% of the data ( $\rho = 1/10$ ):  $r_{10}(1/10) = 0.8$
- Estimate based on 1% of the data ( $\rho = 1/100$ ):  $r_{10}(1/100) = 0.63$

$k$ -nearest neighbors is **not local** in **higher dimensions**!  $\rightarrow$  Far-away data points may not be good predictors for the behavior of the function at  $x$ .

# Probabilistic models for learning

Up to now we assumed

$$y = \tilde{f}(\mathbf{x})$$

Typically assume a **probabilistic model** of the form

$$y = \tilde{f}(\mathbf{x}) + \varepsilon$$

with  $\varepsilon$  **independent** of  $\mathbf{x}$  and  $\mathbb{E}[\varepsilon] = 0$ .

Hence,

$$\mathbb{E}[y|\mathbf{x} = \mathbf{x}] = \tilde{f}(\mathbf{x})$$

How do we effectively use the training data  $\mathcal{D}$  to guide the learning of  $f(\mathbf{x})$ ?

# Probabilistic models for learning

The space of all possible regression functions  $f(\mathbf{x})$  is enormous!

**Idea:** Consider a **parametric form** of  $f(\mathbf{x})$ ,  $f_{\theta}(\mathbf{x})$ , with parameters  $\theta$ .

Example: **Linear regression**,

$$\begin{aligned}\hat{y} &= \sum_{j=0}^p x_j \beta_j + \varepsilon \\ &= \mathbf{x}^T \boldsymbol{\beta} + \varepsilon.\end{aligned}$$

Linear function of the parameters  $\beta_0, \dots, \beta_p$  and of the input variables  $x_1, \dots, x_p$ .



# Probabilistic models for learning

More in general we may consider

$$f_{\theta}(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^M \theta_i \phi_i(\mathbf{x})$$

$\phi_j$ : basis functions or basis expansion

$f_{\theta}(\mathbf{x}, \mathbf{w})$ : linear basis expansion

Resulting model is much richer, but still a linear function in  $\theta$ !

# Maximum likelihood learning

- Assume  $f_{\theta}(x)$  fixed
- We want to learn  $\theta$

Before (least-squares regression):

$$\theta^* = \operatorname{argmin}_{\theta} \sum_{i=1}^N (y_i - f_{\theta}(x_i))^2$$

**Maximum likelihood:** select  $\theta$  for which the training set  $\mathcal{D}$  has the **maximum probability** of being observed.

# Maximum likelihood learning

Choose  $\theta$  that maximizes

$$p(y_{\mathcal{D}} | \mathbf{x}_{\mathcal{D}}, \theta) = \prod_{i=1}^N p(y_i | \mathbf{x}_i, \theta)$$

or, equivalently, the **log-likelihood (LL) function**

$$\ln p(y_{\mathcal{D}} | \mathbf{x}_{\mathcal{D}}, \theta) = \sum_{i=1}^N \ln p(y_i | \mathbf{x}_i, \theta)$$

# Maximum likelihood learning

$$\begin{aligned}\theta_{\text{ML}} &= \operatorname{argmax}_{\theta} \ln p(y_{\mathcal{D}} | \mathbf{x}_{\mathcal{D}}, \theta) \\ &= \operatorname{argmin}_{\theta} - \ln p(y_{\mathcal{D}} | \mathbf{x}_{\mathcal{D}}, \theta) \\ &= \operatorname{argmin}_{\theta} - \sum_{i=1}^N \ln p(y_i | \mathbf{x}_i, \theta) \\ &= \operatorname{argmin}_{\theta} - \frac{1}{N} \sum_{i=1}^N \ln p(y_i | \mathbf{x}_i, \theta)\end{aligned}$$

If  $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ , from  $y = f_{\theta}(\mathbf{x}) + \varepsilon$ :

$$y_i | \mathbf{x}_i, \theta \sim \mathcal{N}(y_i; f_{\theta}(\mathbf{x}_i), \sigma^2)$$

## Maximum likelihood learning

$$\begin{aligned}\boldsymbol{\theta}_{\text{ML}} &= \underset{\boldsymbol{\theta}}{\operatorname{argmin}} -\frac{1}{N} \sum_{i=1}^N \ln p(y_i | \boldsymbol{x}_i, \boldsymbol{\theta}) \\ &= \underset{\boldsymbol{\theta}}{\operatorname{argmin}} -\frac{1}{N} \sum_{i=1}^N \ln \mathcal{N}(y_i; f_{\boldsymbol{\theta}}(\boldsymbol{x}_i), \sigma^2)\end{aligned}$$

# Maximum likelihood learning

# The need of structured regression models

**Goal:** Choose a function  $f \in \mathcal{F}$  that minimizes a given loss function  $L(\hat{y}) = L(f; \mathcal{D})$  based on training set  $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}, i \in [N]$ .

Example:

$$\theta^* = \underset{\theta}{\operatorname{argmin}} \operatorname{RSS} \triangleq \sum_{i=1}^N (y_i - f_{\theta}(\mathbf{x}_i))^2.$$

**Observation:** If  $\mathcal{F}$  set of all possible functions, can make  $\operatorname{RSS}(f(\mathbf{x})) = 0$  (any  $f(\mathbf{x})$  that passes through the training points)

... but not all will **generalize** well to new data.

Need to impose some **constraints** on  $f(\mathbf{x})$ !

# The need of structured regression models

Which **constraints** should we impose?

- Restrict to parametric functions  $f_\theta$  (**linear regression**:  $\mathcal{F}$  family of all linear functions)
- Smoother functions
- ...

Three classes of structured regression models:

- **Roughness penalty**
- **Kernel methods**
- **Basis functions and dictionary methods**



# Class 1: Roughness penalty

Assuming a measure of “niceness” (e.g., **smoothness**)  $J(f)$ ,

$$f(\mathbf{x}) = \operatorname{argmin}_{f \in \mathcal{F}: L(f; \mathcal{D})=0} J(f)$$

**Example:** smoothing splines

For one-dimensional data  $x \in [0, 1]$ ,  $\mathcal{F}$  is the family of all twice-differentiable functions, and we choose  $J(f)$  as

$$J(f) = \int_0^1 (f''(x))^2 dx$$

Can **relax requirement** that  $L(f; \mathcal{D}) = 0$  via

$$f(\mathbf{x}) = \operatorname{argmin}_{f \in \mathcal{F}} L(f; \mathcal{D}) + \lambda J(f)$$

**Regularization methods:** Trade-off between loss and smoothness

## Class 2: Kernel methods

Estimate regression (or classification) function in a **local neighborhood**

- Need to specify nature of local neighborhood and class of functions used for local fit

**Simplest form:** **Nadaraya-Watson** weighted average,

$$f(\mathbf{x}) = \frac{\sum_{i=1}^N K_{\lambda}(\mathbf{x}, \mathbf{x}_i) y_i}{\sum_{i=1}^N K_{\lambda}(\mathbf{x}, \mathbf{x}_i)}$$

$K_{\lambda}(\mathbf{x}, \mathbf{x}_i)$ : **Kernel function**; assigns weights to  $\mathbf{x}_i$  depending on closeness to  $\mathbf{x}$

**Example 1:** **k-nearest neighbors**

$$K_k(\mathbf{x}, \mathbf{x}_i) = \mathbb{1}\{\|\mathbf{x}_i - \mathbf{x}\| \leq \|\mathbf{x}_{(k)} - \mathbf{x}\|\}$$

$\mathbf{x}_{(k)}$ :  $k$ -th closest input in data set to  $\mathbf{x}$

**Example 2:** **Gaussian kernel**

$$K_{\lambda}(\mathbf{x}, \mathbf{x}_i) = \frac{1}{\lambda} \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{2\lambda}\right)$$

## Class 2: Kernel methods

More in general,

$$\boldsymbol{\theta}^* = \underset{\boldsymbol{\theta}}{\operatorname{argmin}} \sum_{i=1}^N K_{\lambda}(\boldsymbol{x}, \boldsymbol{x}_i) (y_i - f_{\boldsymbol{\theta}}(\boldsymbol{x}_i))^2$$

## Class 3: Basis functions and dictionary methods

$f$  modeled as a linear expansion of basis functions:

$$f_{\theta}(\mathbf{x}, \mathbf{w}) = \sum_{i=1}^M \theta_i \phi_i(\mathbf{x})$$

# The bias–variance trade-off

- $k$ -NN method:  $k$ , the number of nearest neighbors
- Linear models:  $M$ , the model order (the number of basis functions)
- Roughness penalty methods:  $\lambda$ , the weight of the penalty term

Parameters control the capacity to fit data

Higher capacity (higher complexity)  $\longrightarrow$  fit training data more accurately  
... but unlikely to generalize well

Low capacity  $\longrightarrow$  cannot capture all variations present on data and may generalize poorly

# The bias–variance trade-off

**Idea:** Divide expected prediction error into its components (**bias** and **variance** for the squared error loss; **approximation** and **estimation** error for general case).

We consider:

$$y = \tilde{f}(\mathbf{x}) + \varepsilon,$$

with

- $\varepsilon \sim \mathcal{N}(0, \sigma^2) \longrightarrow \mathbb{E}[y|\mathbf{x}] = \tilde{f}(\mathbf{x})$
- $\mathbf{x} = \mathbf{x}$  fixed
- square loss function  $\ell(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$
- Large number of datasets drawn from  $p(y, \mathbf{x})$

**Expected prediction (generalization) error:**

$$\text{Err}(y, f(\mathbf{x})) = \mathbb{E}_{\mathcal{D}, \mathbf{x}, y}[\ell(y, f(\mathbf{x}))]$$

## The bias–variance trade-off

$$\text{Err}(y, f(\mathbf{x})) = \mathbb{E}_{\mathcal{D}, \mathbf{x}, y}[\ell(y, f(\mathbf{x}))]$$

For  $\mathbf{x} = \mathbf{x}$  fixed:

$$\begin{aligned}\text{Err}(\mathbf{x}) &= \mathbb{E}_{y|\mathbf{x}, \mathcal{D}}[\ell(y, f(\mathbf{x})) | \mathbf{x} = \mathbf{x}] \\ &= \mathbb{E}_{y|\mathbf{x}} \mathbb{E}_{\mathcal{D}}[\ell(y, f(\mathbf{x})) | \mathbf{x} = \mathbf{x}] \\ &= \mathbb{E}_{y|\mathbf{x}} \mathbb{E}_{\mathcal{D}}[(y - f(\mathbf{x}))^2 | \mathbf{x}] \\ &= \mathbb{E}_{y|\mathbf{x}} \mathbb{E}_{\mathcal{D}}[(y - \mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] + \mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x})] + \mathbb{E}_{\mathcal{D}}[f(\mathbf{x})] - f(\mathbf{x}))^2] \\ &= \mathbb{E}_{y|\mathbf{x}}[(y - \mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}])^2] + (\mathbb{E}_{y|\mathbf{x}}[y|\mathbf{x}] - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x})])^2 + \mathbb{E}_{y|\mathbf{x}} \mathbb{E}_{\mathcal{D}}[(\mathbb{E}_{\mathcal{D}}[f(\mathbf{x})] - f(\mathbf{x}))^2] \\ &= \mathbb{E}_{y|\mathbf{x}}[(y - \tilde{f}(\mathbf{x}))^2] + (\tilde{f}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x})])^2 + \mathbb{E}_{y|\mathbf{x}} \mathbb{E}_{\mathcal{D}}[(\mathbb{E}_{\mathcal{D}}[f(\mathbf{x})] - f(\mathbf{x}))^2] \\ &= \text{Var}[y|\mathbf{x}] + (\tilde{f}(\mathbf{x}) - \mathbb{E}_{\mathcal{D}}[f(\mathbf{x})])^2 + \mathbb{E}_{\mathcal{D}}[(\mathbb{E}_{\mathcal{D}}[f(\mathbf{x})] - f(\mathbf{x}))^2]\end{aligned}$$

# The bias–variance trade-off


$$\begin{aligned}\text{Err}(\boldsymbol{x}) &= \text{Var}[y|\boldsymbol{x}] + (\tilde{f}(\boldsymbol{x}) - \mathbb{E}_{\mathcal{D}}[f(\boldsymbol{x})])^2 + \mathbb{E}_{\mathcal{D}} [(\mathbb{E}_{\mathcal{D}}[f(\boldsymbol{x})] - f(\boldsymbol{x}))^2] \\ &= \sigma^2 + \underbrace{(\tilde{f}(\boldsymbol{x}) - \mathbb{E}_{\mathcal{D}}[f(\boldsymbol{x})])^2}_{\text{bias}} + \text{Var}_{\mathcal{D}} [f(\boldsymbol{x})]\end{aligned}$$

## Bias-variance decomposition

- $\sigma^2$ : **Irreducible error** due to data randomness
- **Bias**: **Approximation error** due to the limited flexibility model
- **Variance**: **Estimation error**; sensitivity/variability of model due to randomness in  $\mathcal{D}$



# The bias–variance trade-off



Bias-variance trade-off:

- High capacity models: low bias, high variance (overfitting)
- Low capacity models: high bias, low variance (underfitting)

# The bias–variance trade-off or least squares regression

$$\text{Err}(\boldsymbol{x}) = \sigma^2 + \sigma^2 \boldsymbol{x}^\top \mathbb{E}_{\mathcal{D}} \left[ (\mathbf{X}^\top \mathbf{X})^{-1} \right] \boldsymbol{x}$$

and

$$\begin{aligned} \text{Err}(y, f(\boldsymbol{x})) &= \mathbb{E}_{\mathbf{x}}[\text{Err}(\mathbf{x})] \\ &= \sigma^2 + \sigma^2 \mathbb{E}_{\mathbf{x}, \mathcal{D}} \left[ \mathbf{x}^\top (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{x} \right] \end{aligned}$$

Tutorial exercise!

# Reading

“The elements of statistical learning,”  
Chapters 1 and 2